

Machine Learning Project work: “Real-Time Fire Detection on the Edge”

Mario Vento, Pasquale Foggia
and Diego Gragnaniello

SEPTEMBER
11-15.2023
UDINE
ITALY

ICIAP
2023
UDINE

ONFIRE International Contest!!!

- ◆ This very same contest has been accepted by the **22nd International Conference on Image Analysis and Processing** (Udine, September 11-15)
- ◆ International research teams will participate (submit their technique) and present their method
- ◆ The **student team ranking first will be invited to participate in the conference** to present their solution

The task to address

- ◆ The aim of this project is to detect FIRE in videos acquired by video surveillance cameras
 - ◆ Static camera
 - ◆ Classify a sequence of frames (images) as either (P)ositive (fire) or (N)egative (no fire)
 - ◆ Fire can occur at a specific frame that must be detected

N



P



The task to address

- ◆ Videos are acquired by different cameras at different resolutions/frame rate
- ◆ May contain over-imposed text

N



P



The task to address

- ◆ Videos are acquired during daylight/night, indoor/outdoor and by different distances
- ◆ Note: at a distance, the smoke only may be visible

N



The task to address

- ◆ Fire may not be visible in the first frames
- ◆ You have to provide the first frame in which it is visible



The task to address

- ◆ Videos may contain moving objects



The task to address

- ◆ Videos may contain Fire-like (yellowish-reddish) or smoke-like (fog, clouds) objects



The training set

- ◆ 330 videos:
 - ◆ 246 FIRE
 - ◆ 84 NO FIRE
- ◆ We will provide the link to online datasets from which these are selected
- ◆ You can expand it with more data of your choice

IMPORTANT: Model training and validation is not a feed-forward process

- ◆ Alternate trainings and validations:
 - ◆ After each validation, try to understand which samples are misclassified and why
 - ◆ Change your model and/or training set to improve the performance

IMPORTANT: Model training and validation is not a feed-forward process

- ◆ Alternate trainings and validations:
 - ◆ If the performance seems very good, be sure that the validation set is challenging enough
 - ◆ Keep track of your countermeasures/improvements for the final project presentation

What you can use

- ◆ You can extend the provided training set
- ◆ You can use data augmentation techniques
- ◆ Any algorithm (whatever kind of classifier, including non-neural ones, preprocessing, training strategy, validation)
 - But you must be able to explain what you have used
 - It must be **runnable inside Google Colab**
- ◆ You can use local computing power (you are not limited to Colab) for the training

Model evaluation

We define:

- ◆ g_i the first frame in which the fire is visible
 - ◆ p_i the first frame in which the fire is detected
 - ◆ Δt a guard time equals 5 seconds
-
- ◆ True Positive (TP): all the detections in (P)ositive videos for which $p \geq \max(0, g - \Delta t)$
 - ◆ False Positive (FP): all the detections occurring at any time in (N)egative videos or in (P)ositive videos for which $p < \max(0, g - \Delta t)$
 - ◆ False Negative (FN): the set of positive videos for which no fire detection occurs.

Model evaluation: the contributions to the final score

◆ The final score is obtained by combining two different contributions:

1. Detection performance metrics:

- ◆ Precision
- ◆ Recall
- ◆ The delay between the fire ignition and detection

2. Complexity:

- ◆ The processing capabilities in terms of frame rate
- ◆ The GPU memory required

Model evaluation: the contributions to the final score

◆ The final score is obtained by combining:

◆ Precision

$$P = \frac{|TP|}{|TP| + |FP|}$$

◆ Recall

$$R = \frac{|TP|}{|TP| + |FN|}$$

Model evaluation: the contributions to the final score

◆ The final score is obtained by combining:

◆ Precision

$$P = \frac{|TP|}{|TP| + |FP|}$$

◆ Recall

$$R = \frac{|TP|}{|TP| + |FN|}$$

◆ Detection delay

$$D = \frac{\sum_{i=1}^{|TP|} d_i}{|TP|}, \quad d_i = |p_i - g_i|, \quad D_n = \frac{\max(0; 60 - D)}{60}$$

Max delay:
60 seconds

Time instant
when fire is
predicted

Time instant
when fire is
visible

Model evaluation: the contributions to the final score

◆ The final score is obtained by combining:

◆ Precision

$$P = \frac{|TP|}{|TP| + |FP|}$$

◆ Recall

$$R = \frac{|TP|}{|TP| + |FN|}$$

◆ Detection delay

$$D = \frac{\sum_{i=1}^{|TP|} d_i}{|TP|}, \quad d_i = |p_i - g_i|, \quad D_n = \frac{\max(0; 60 - D)}{60}$$

◆ Processing frame rate

$$PFR = \frac{1}{\frac{\sum_{i=1}^{|N|} t_i}{|N|}}, \quad PFR_{delta} = \max\left(0; \frac{PFR_{target}}{PFR} - 1\right)$$

◆ Memory occupancy

$$MEM_{delta} = \max\left(0; \frac{MEM}{MEM_{target}} - 1\right)$$

Penalize:

memory
occupancy
greater than
 MEM_{target}

proc. frame
rate lower
than
 PFR_{target}

Model evaluation: the contributions to the final score

◆ The final score is obtained by combining:

◆ Precision

$$P = \frac{|TP|}{|TP| + |FP|}$$
$$R = \frac{|TP|}{|TP| + |FN|}$$

◆ Recall

◆ Detection delay

$$D = \frac{\sum_{i=1}^{|TP|} d_i}{|TP|}, \quad d_i = |p_i - g_i|, \quad D_n = \frac{\max(0; 60 - D)}{60}$$

◆ Processing frame rate

$$PFR = \frac{1}{\frac{\sum_{i=1}^{|N|} t_i}{|N|}}, \quad PFR_{delta} = \max\left(0; \frac{PFR_{target}}{PFR} - 1\right)$$

◆ Memory occupancy

$$MEM_{delta} = \max\left(0; \frac{MEM}{MEM_{target}} - 1\right)$$

◆ The final score

$$FDS = \frac{P * R * D_n}{(1 + PFR_{delta}) * (1 + MEM_{delta})}$$

Share the load

- ◆ Each member of the team will be requested to submit an estimate of the individual effort contributed by all members
 - To prevent "free riders"
 - Submissions will be "blind" (each member will not see the submissions of other members)

IMPORTANT: Don't forget to

- ◆ Write the **names of all the team members** in the Google Drive folder (in a text file)
- ◆ Ensure that the **link** you submit is **readable to anyone** (no authorization must be requested)
- ◆ Make sure that the **test script** is **compliant with the specification** (if you have doubts about the specification, **ask**)